

# Convolutional NNs and Generative Models

Friday  
08h00-09h00

Géraldine Schaller, Matthew Vowels, Bern Winter School on Machine Learning 2024, Muerren

TO COMPLETE YOUR REGISTRATION, PLEASE TELL US WHETHER OR NOT THIS IMAGE CONTAINS A STOP SIGN:



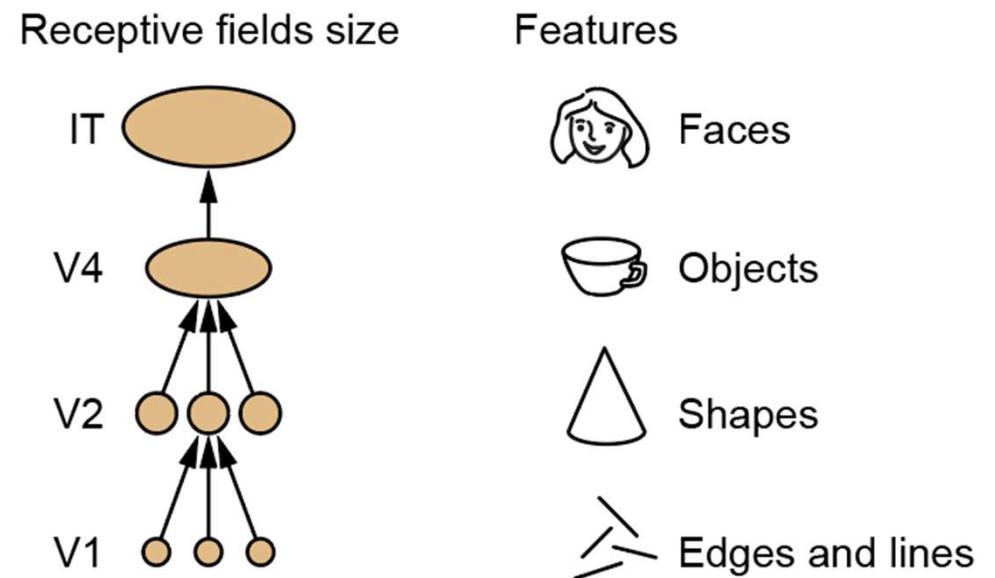
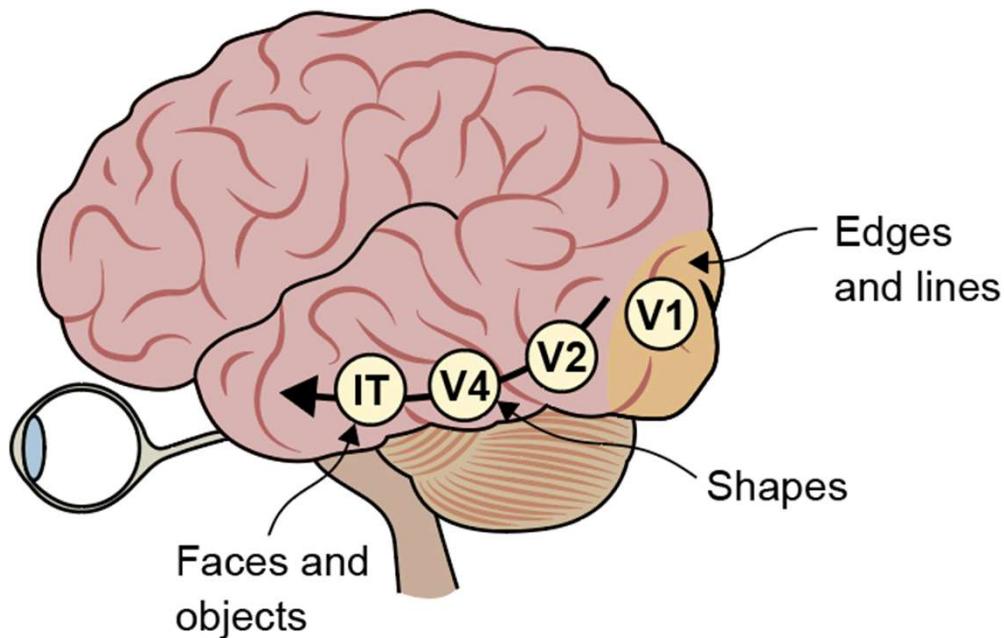
NO

YES

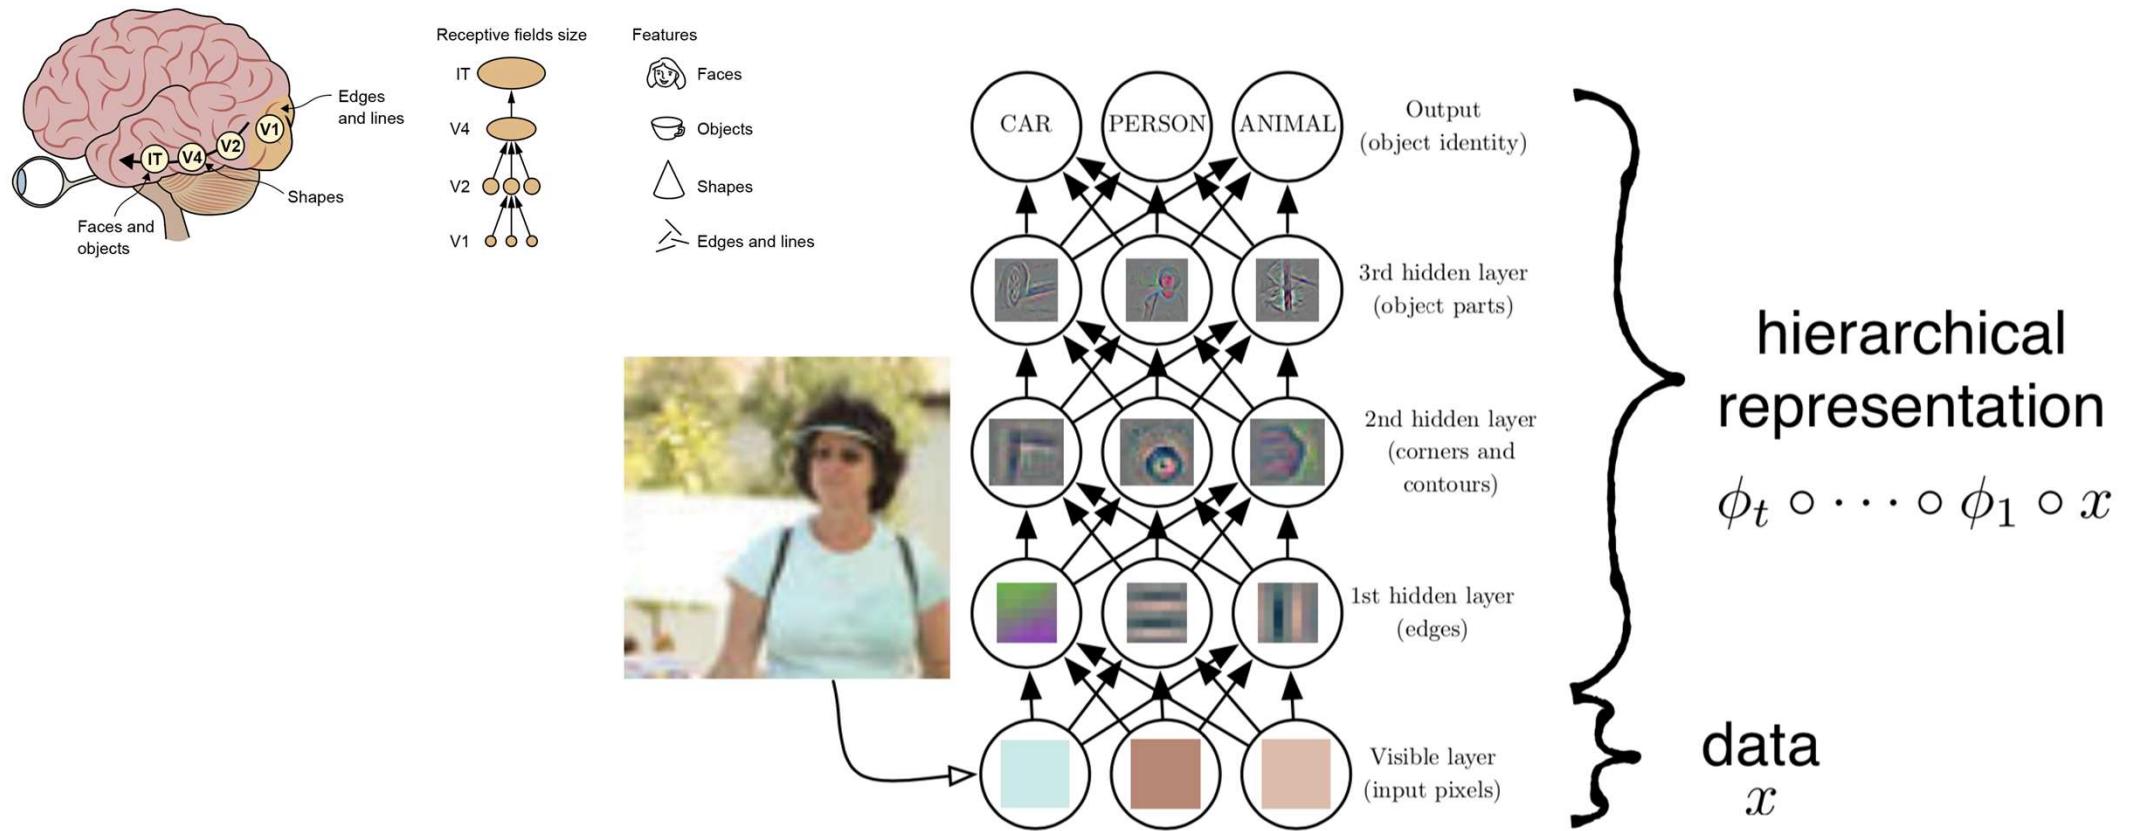
ANSWER QUICKLY—OUR SELF-DRIVING CAR IS ALMOST AT THE INTERSECTION.

SO MUCH OF "AI" IS JUST FIGURING OUT WAYS TO OFFLOAD WORK ONTO RANDOM STRANGERS.

# Human vision



# Computer vision inspired from Nature

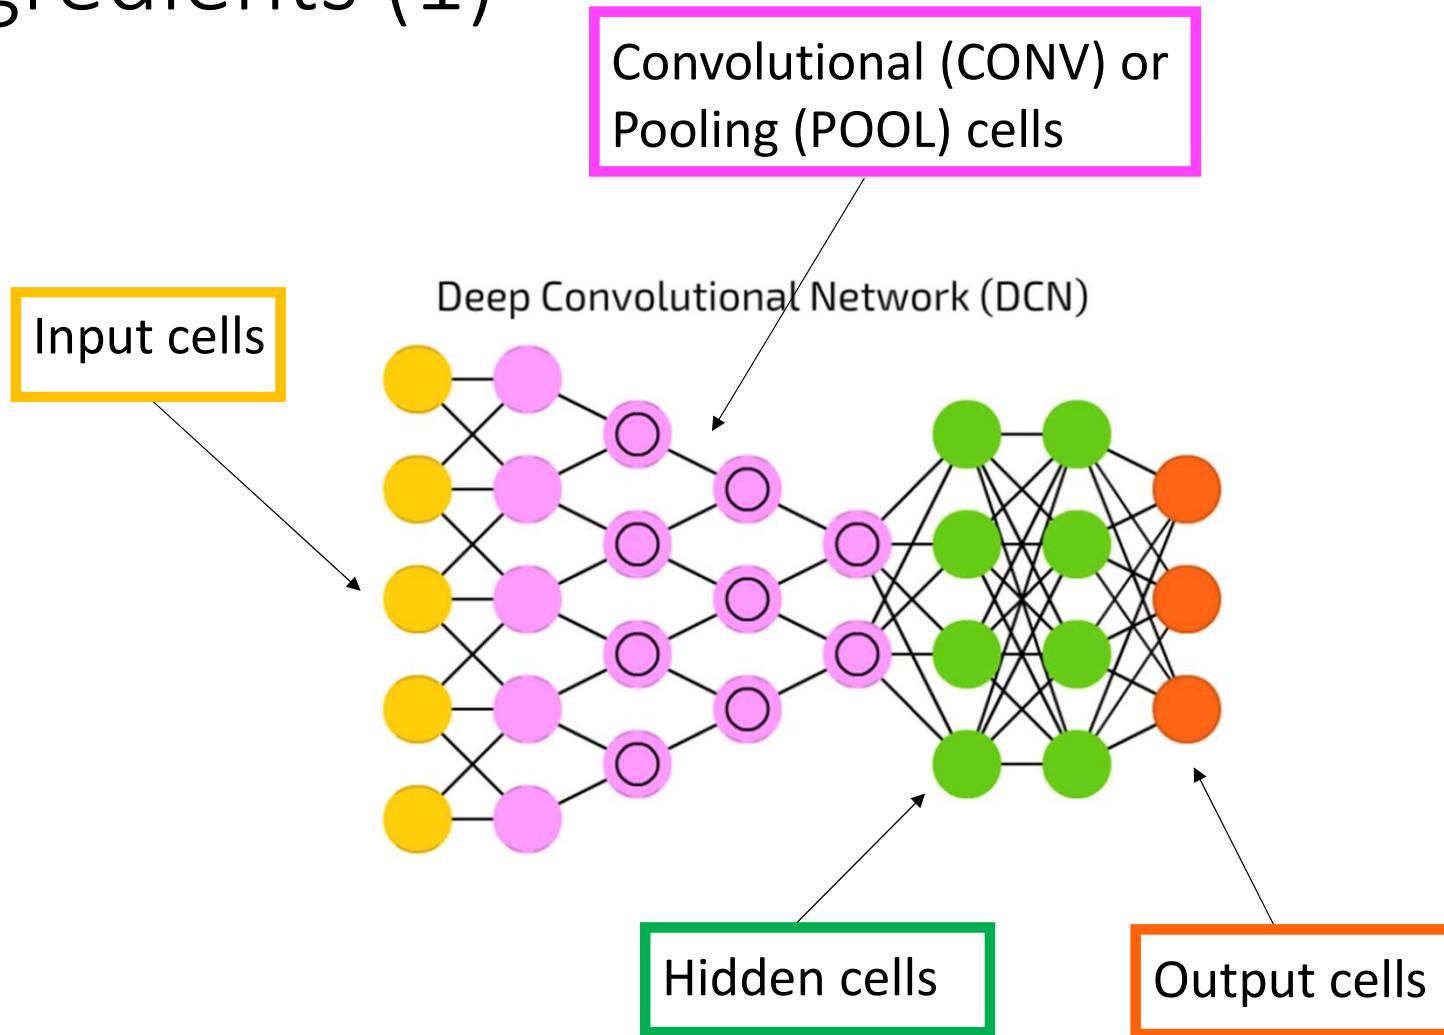


# Input data

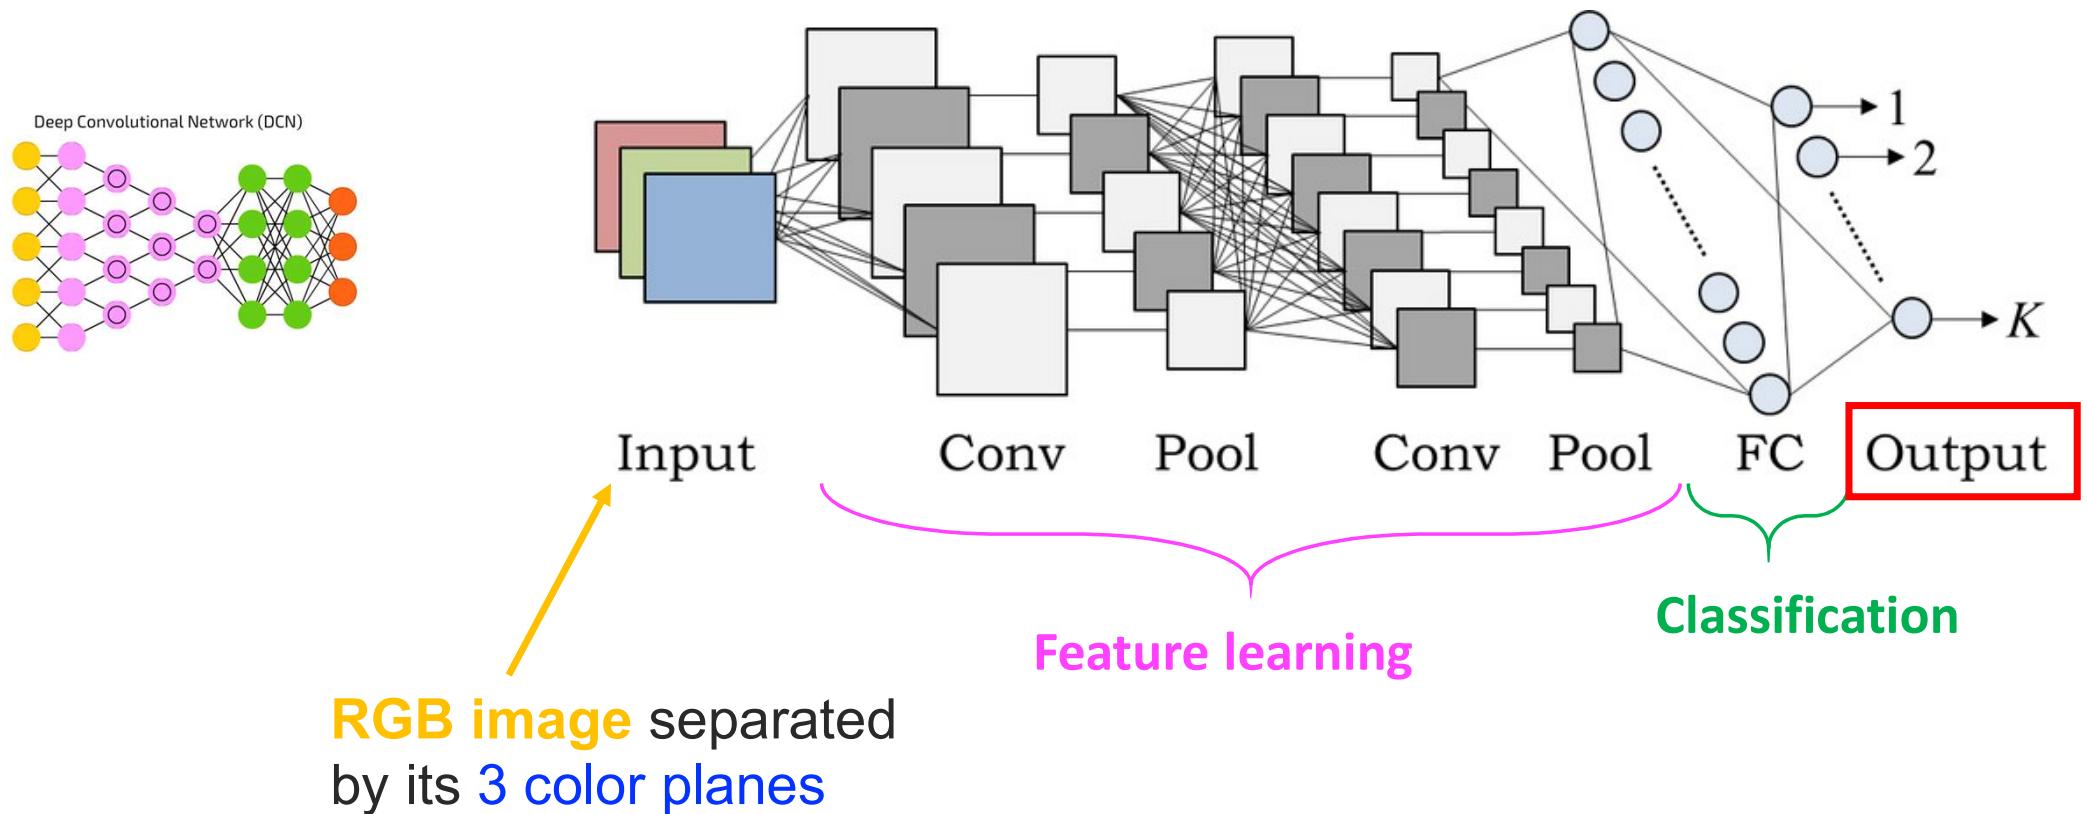
- Specialized Neural Network for data arranged **on a grid**
  - Images
  - DNA sequences
  - ...

	A	C	G	T	W	S	M	K	R	Y	B	D	H	V	N	Z
A	1	0	0	0	1/2	0	1/2	0	1/2	0	0	1/3	1/3	1/3	1/4	0
C	0	1	0	0	0	1/2	1/2	0	0	1/2	1/3	0	1/3	1/3	1/4	0
G	0	0	1	0	0	1/2	0	1/2	1/2	0	1/3	1/3	0	1/3	1/4	0
T	0	0	0	1	1/2	0	0	1/2	0	1/2	1/3	1/3	1/3	0	1/4	0

# CNN ingredients (1)



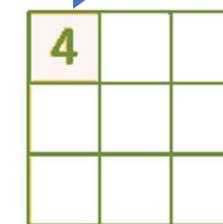
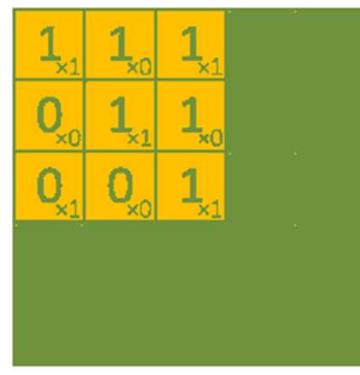
## CNN ingredients (2)



# CONV layer : the Kernel

- Used to **detect features** (vertical/horizontal filter,...)

$$\begin{matrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{matrix}$$



how well the pattern in the kernel matches the content in that part of the image

Convolved Feature

= Feature map

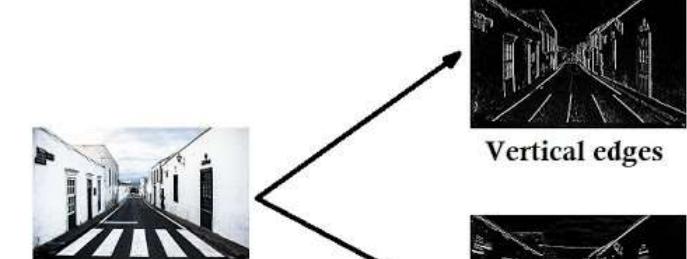
- Different kernels to create different feature maps → learn to see various patterns and details in images

$$\begin{matrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{matrix}$$

Vertical

$$\begin{matrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{matrix}$$

Horizontal



Vertical edges



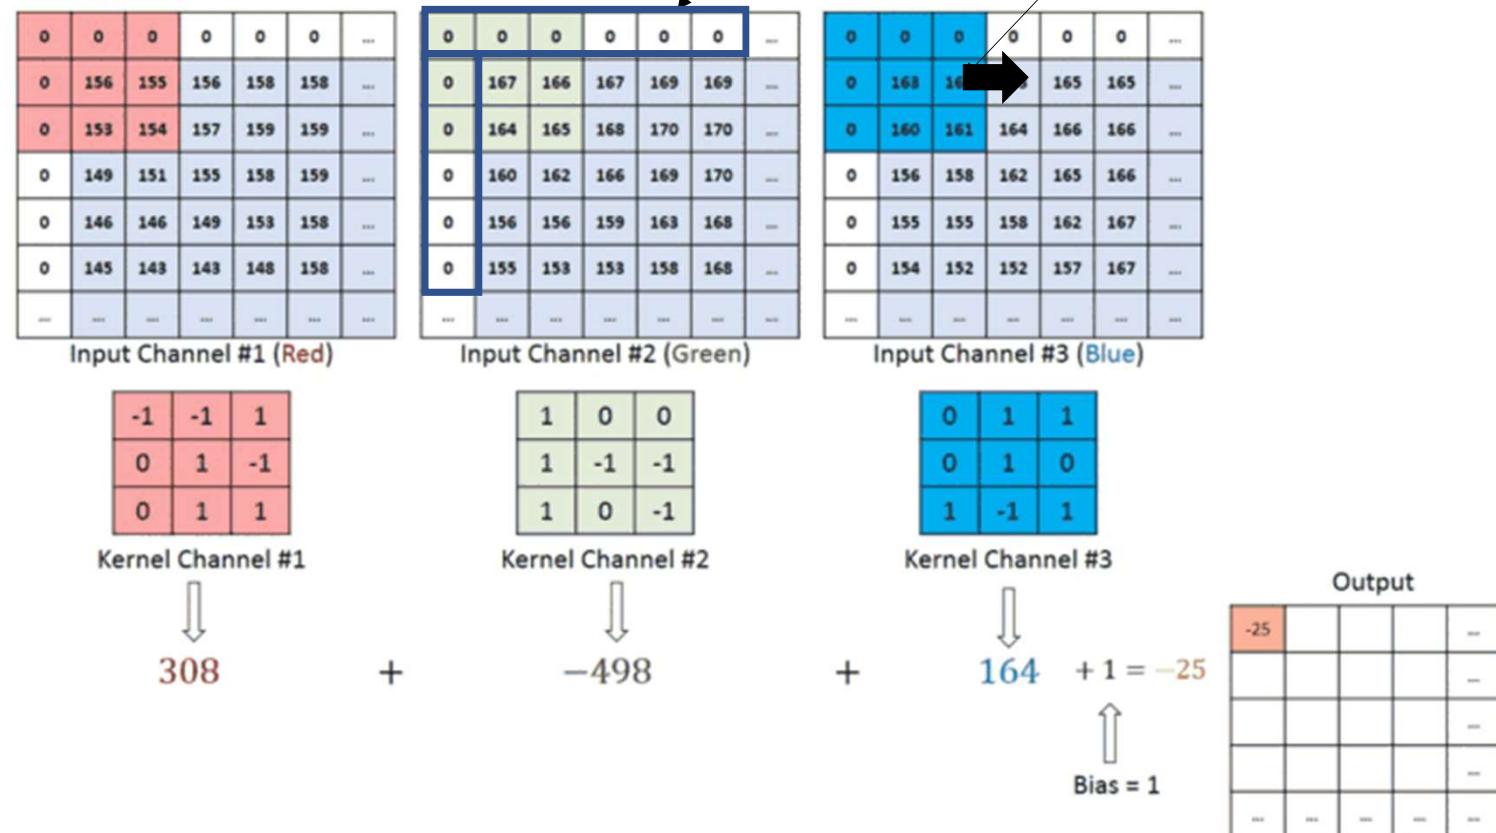
Horizontal edges

# Convolution on RGB image

- In this example, use of a 3x3x3 filter

Width, height, #input channels

In a standard color image, the number of input channels is 3 (for red, green, and blue)



# Kernels (filters)

- Values inside a kernel are learned during the training process of the CNN
- Two aspects to consider for kernels :
  - Architecture :
    - how large they should be
    - How many input channels the kernel operates on
    - how many of them are applied to the input data

# parameters = (Filter Width \* Filter Height \* Input Channels \* Number of Filters / Stride) + Number of Filters

- Initialization : how they should look at the start of training
  - Random initialization (Xavier/He initialization)
  - Transfer learning

# Padding

- Adds zeros around the border of an image



=

Blue	123	94	83	4	30
Red	123	94	83	2	92
Green	34	44	187	92	4
	34	75	232	124	4
	67	83	164	202	



0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

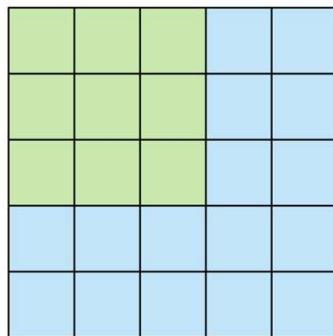
- Use cases :

- Keeps more information at the border of an image
- Allows to use a CONV layer without shrinking the height and width of the volumes
  - *important for deeper networks*

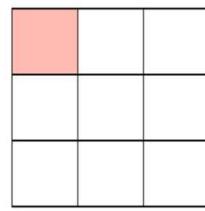
A 3x3 grid of numbers representing an input volume. The central 3x3 area contains values from 123 to 202. This is surrounded by a single layer of zeros labeled "pad". A 2x2 kernel is applied to the input, resulting in a 2x2 output volume where each element is the sum of the 3x3 receptive field. Arrows point from the 3x3 input matrix to this diagram.

# Stride

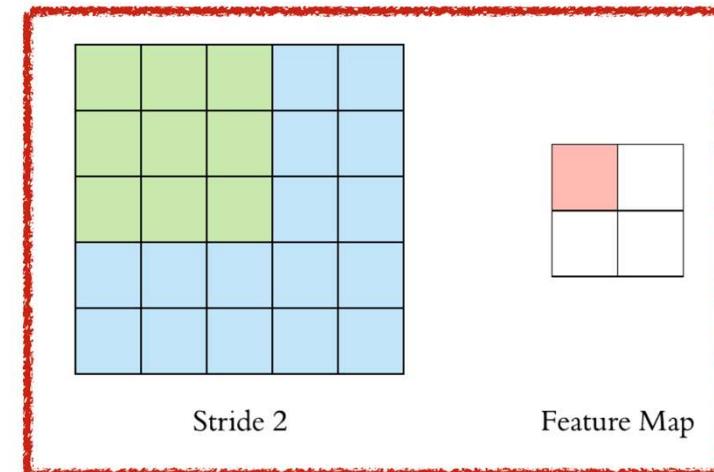
- Can be thought of as a "step size": i.e., **how much you move the filter**
- **Use cases :**
  - Connect a large input layer to a much smaller layer by spacing out the receptive fields (reduce dimensionality)



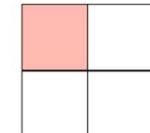
Stride 1



Feature Map



Stride 2



Feature Map

**increasing stride from 1 to 2**



- 300x300 RGB image, how many parameters does a hidden layer have (incl. bias) in the following cases:

- a) Hidden layer with 100 neurons, each fully connected to the input (*no convolution*)

params = (number of input connections + 1) \* number of neurons

ANSWER:

```
>>> params = (300*300*3+1)*100  
>>> print(params)  
27000100
```

- a) Hidden *convolution* layer with 100 filters that are each 5x5

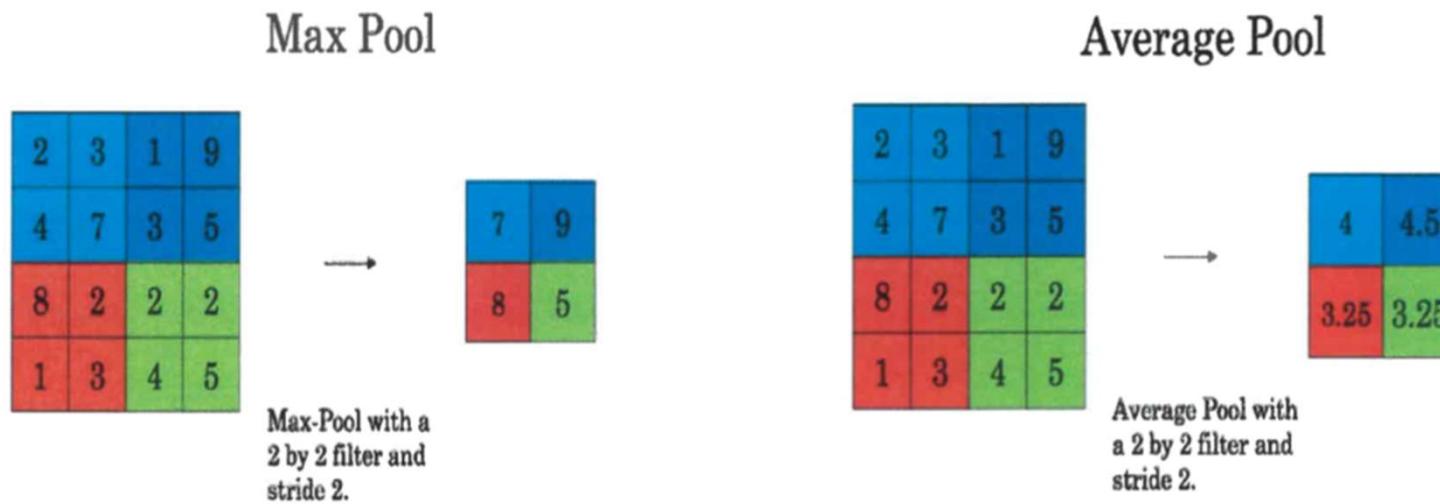
params = (Filter Width \* Filter Height \* Number of Input Channels / Stride<sup>2</sup>)  
\* Number of Filters + Number of Filters

ANSWER:

```
>>> paramsC = (5*5*3)*100+100  
>>> print(paramsC)  
7600
```

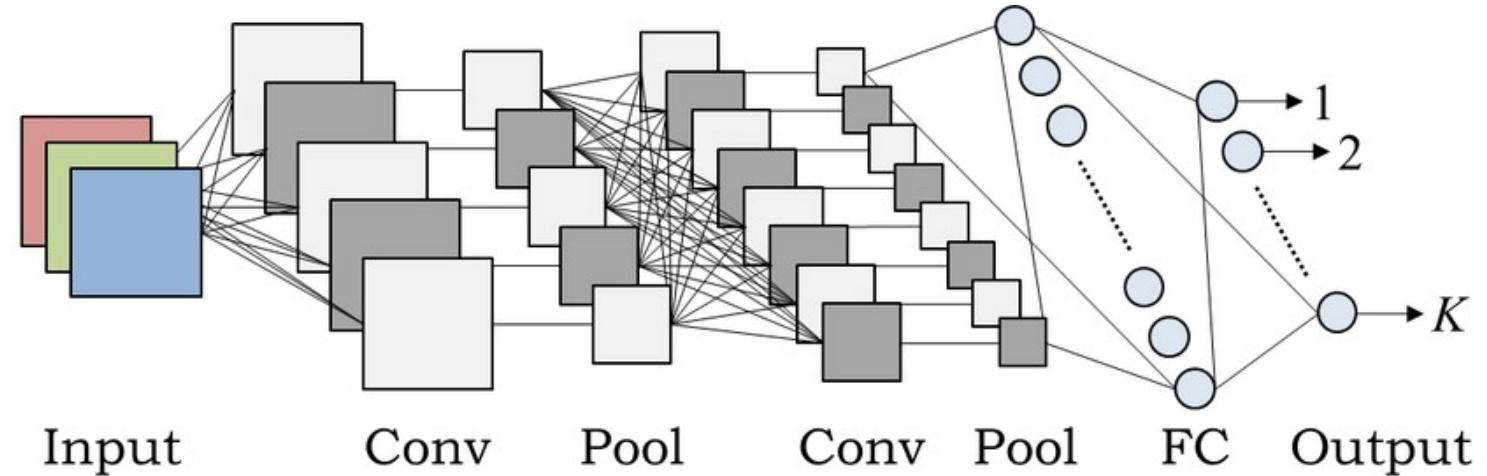
# POOL layer

- Goal is to **subsample** the input image : the exact location of a feature is less important than its rough location relative to other features.
  - Reduce the computational load, memory usage, number of parameters (hence overfitting)



- Pooling neuron has **no weights**, it aggregates the inputs with an **aggregation function (max, mean)**

# CNN Architecture : Summary



detects patterns in  
multiple sub-regions  
in the input field  
using receptive fields

progressively reduces the  
spatial size of the  
representation (fewer  
parameters, less overfitting)

Image gets **smaller and smaller**, but also **deeper and deeper**  
(with more feature maps due to the CONV layers)



## **One Look Is Worth A Thousand Words--**

One look at our line of Republic, Firestone, Miller and United States tires can tell you more than a hundred personal letters or advertisements.

**WE WILL PROVE THEIR VALUE  
BEFORE YOU INVEST ONE DOLLAR  
IN THEM.**

Ever consider buying Supplies from a catalog?

What's the use! Call and see what you are buying. One look at our display of automobile and motorcycle accessories will convince you of the fact.

**THAT WE HAVE EVERYTHING FOR  
THE AUTO**

### **Piqua Auto Supply House**

133 N. Main St.—Piqua, O.

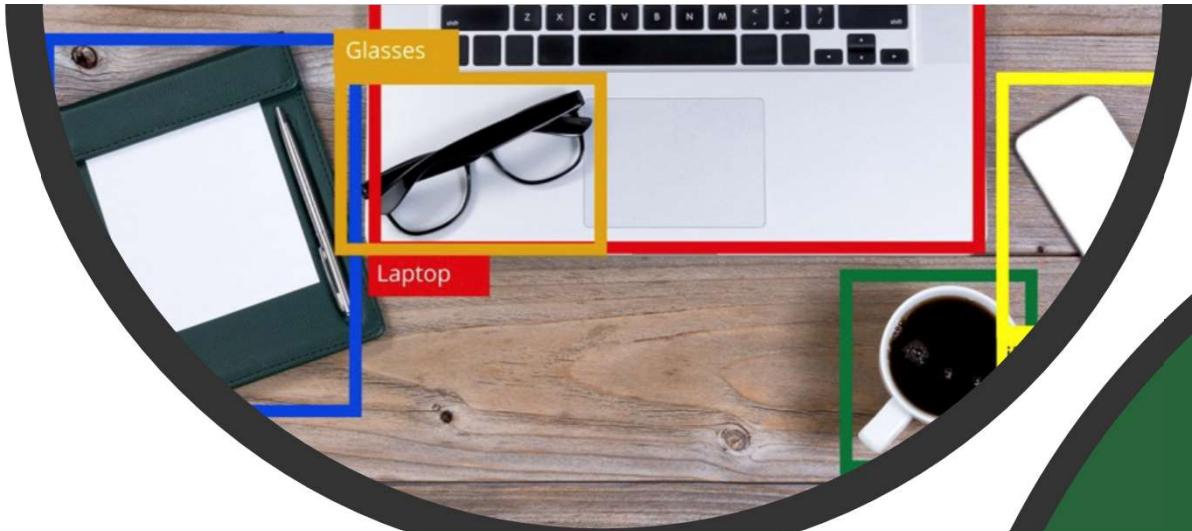


[https://www.youtube.com/watch?v=1zvohULpe\\_0](https://www.youtube.com/watch?v=1zvohULpe_0)

# Object Detection



18



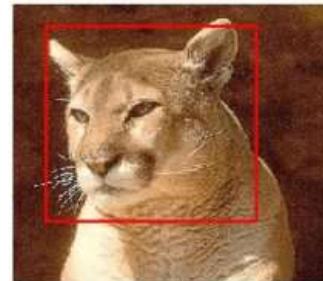
# Introduction

## Classification



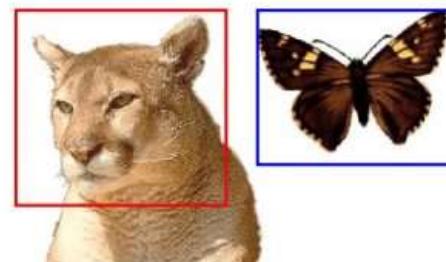
Cougar

## Classification + Localization



Cougar

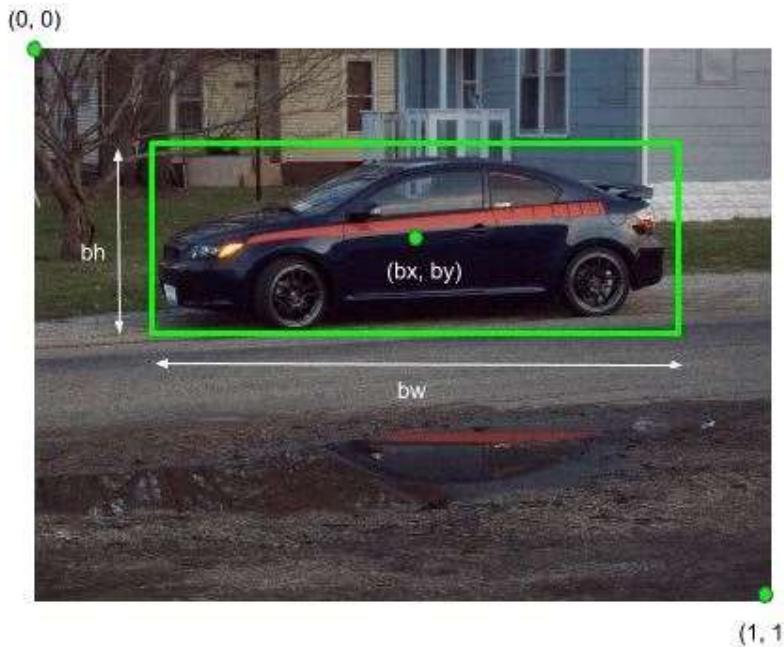
## Object Detection



Cougar, Butterfly

Multiple Objects

# Object localisation



**Target variable :**

$$y = \begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \\ c_4 \end{bmatrix}$$

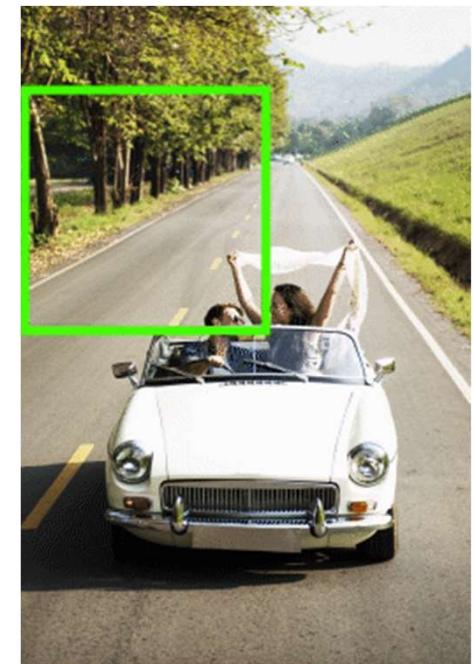
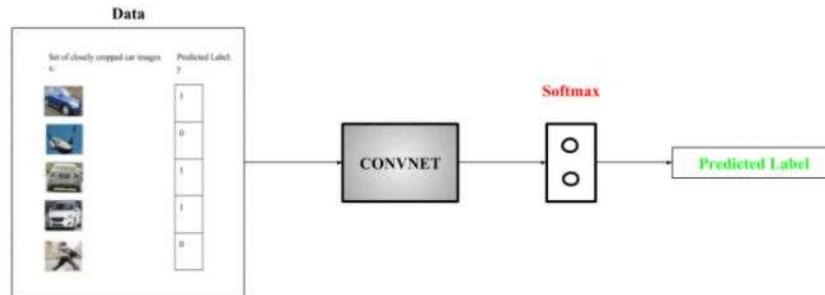
pc = Probability/confidence of an object (i.e the four classes) being present in the bounding box

ci = Probability of the ith class the object belongs to

- **Bounding box parameters :**
  - bx, by : coordinates of the center of the bounding box
  - bw : width of the bounding box w.r.t the image width
  - bh : height of the bounding box w.r.t the image height

# Object detection principle

- 1) A classifier that can classify **closely cropped images** of an object



- 2) Use a **sliding window mechanism** and crop a part of the image in each slide

- 3) Each cropped image is then passed to a **ConvNet** model, which in turn predicts the probability that the cropped image is a car

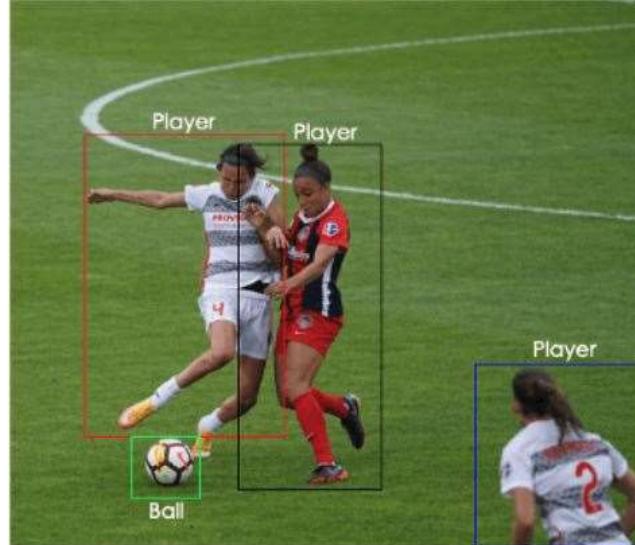
# Self-driving cars

Object Detection Example



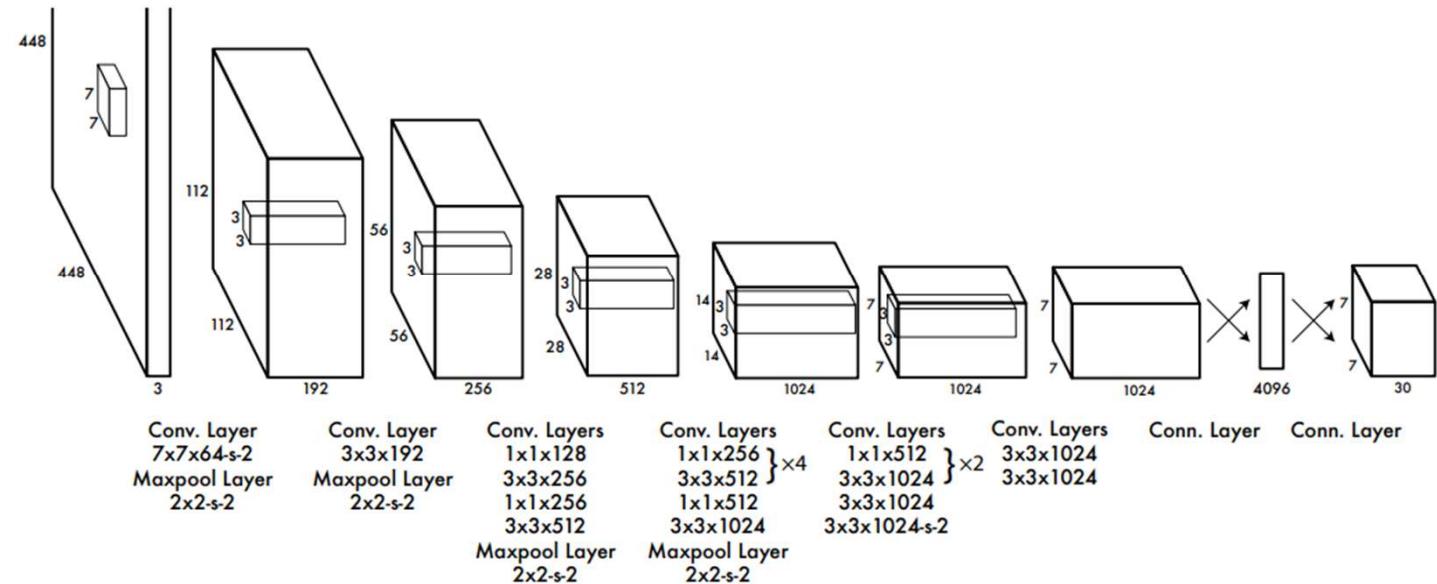
# YOLO algorithm

- You Only Look Once (YOLO) : real-time object detection algorithm
- object detection problem framed as a regression problem instead of a classification task
  - spatially separate bounding boxes and associate probabilities to each of the detected images using a CNN



# YOLO Architecture

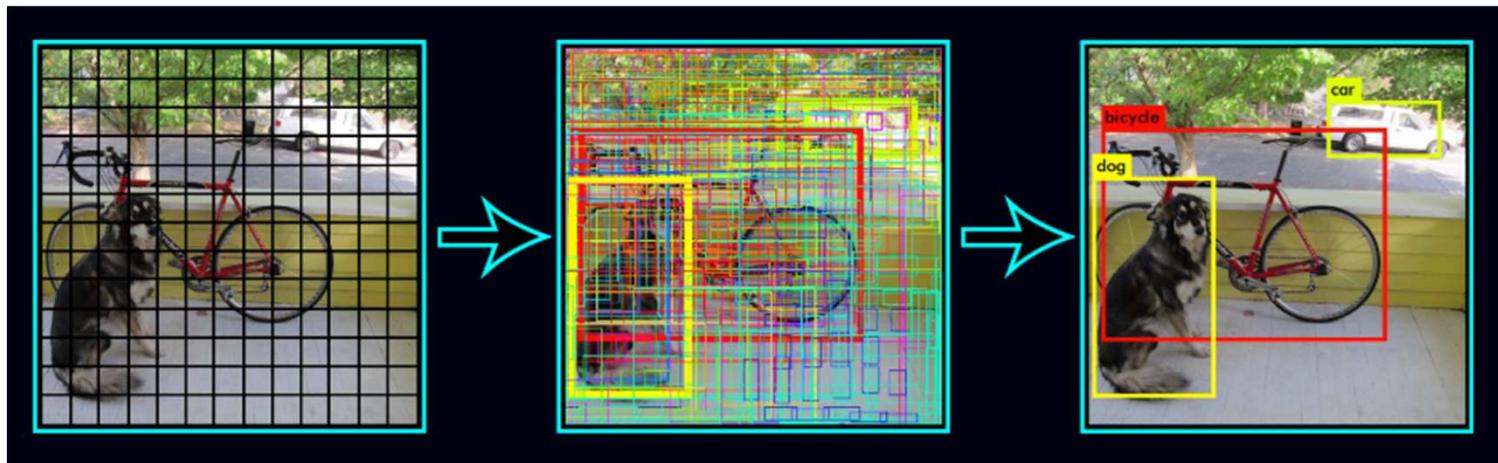
- YOLO architecture is similar to GoogleNet
  - 24 convolutional layers
  - 4 max-pooling layers
  - 2 fully connected layers



- Pretrained **YOLOv2 model**, trained on the COCO image dataset with **classes similar to the Berkeley Driving Dataset**
  - load pretrained weights
  - freeze all the weights except for the last layer during training

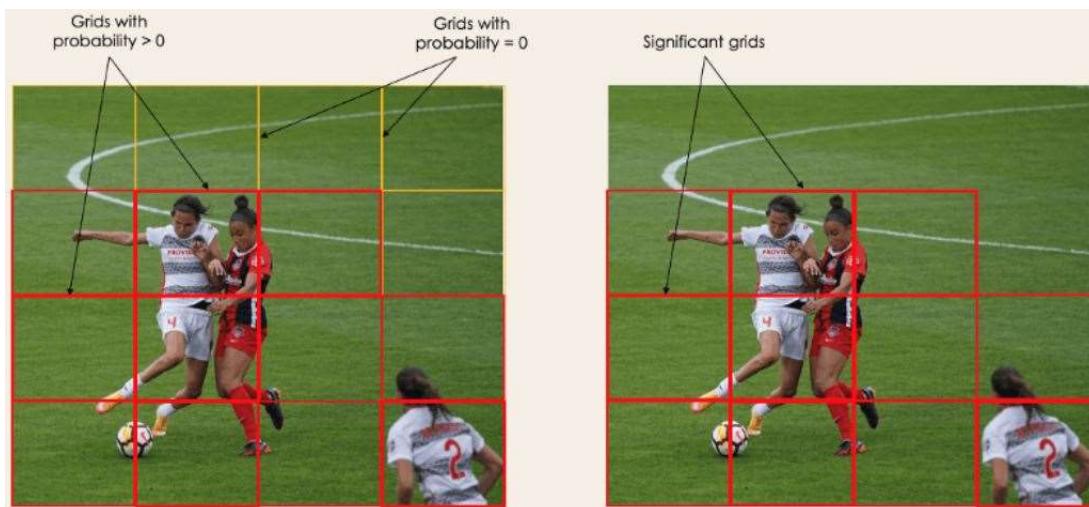
# How does it work

- The algorithm works based on the **four ingredients** :
  - **Residual blocks** : divide the image into NxN grid cells of equal shape
  - **Bounding box regression**
  - **Intersection Over Unions (IoU)**
  - **Non-Max Suppression** : keep only the boxes with the highest probability score of detection.



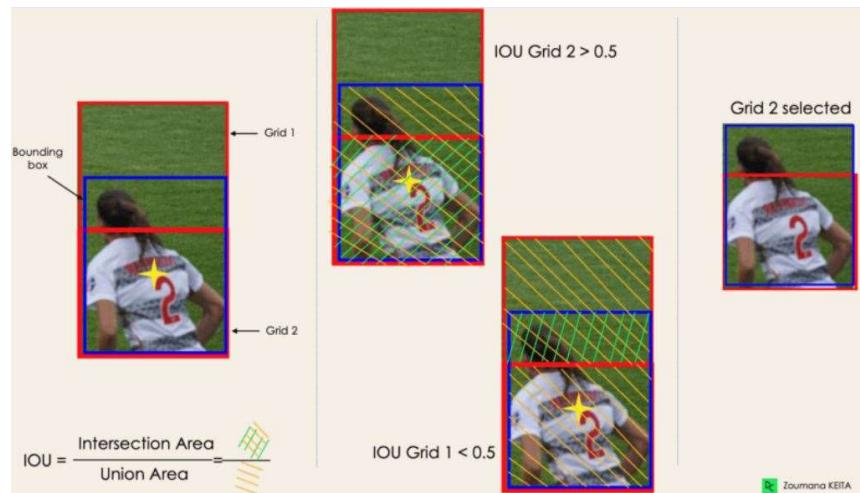
# Bounding box regression

- YOLO determines the attributes of these bounding boxes using a single regression module
  - $Y$  is the final vector representation for each bounding box :  
$$Y = [pc, bx, by, bh, bw, c1, c2]$$
  - $pc$  corresponds to the probability score of the grid containing an object

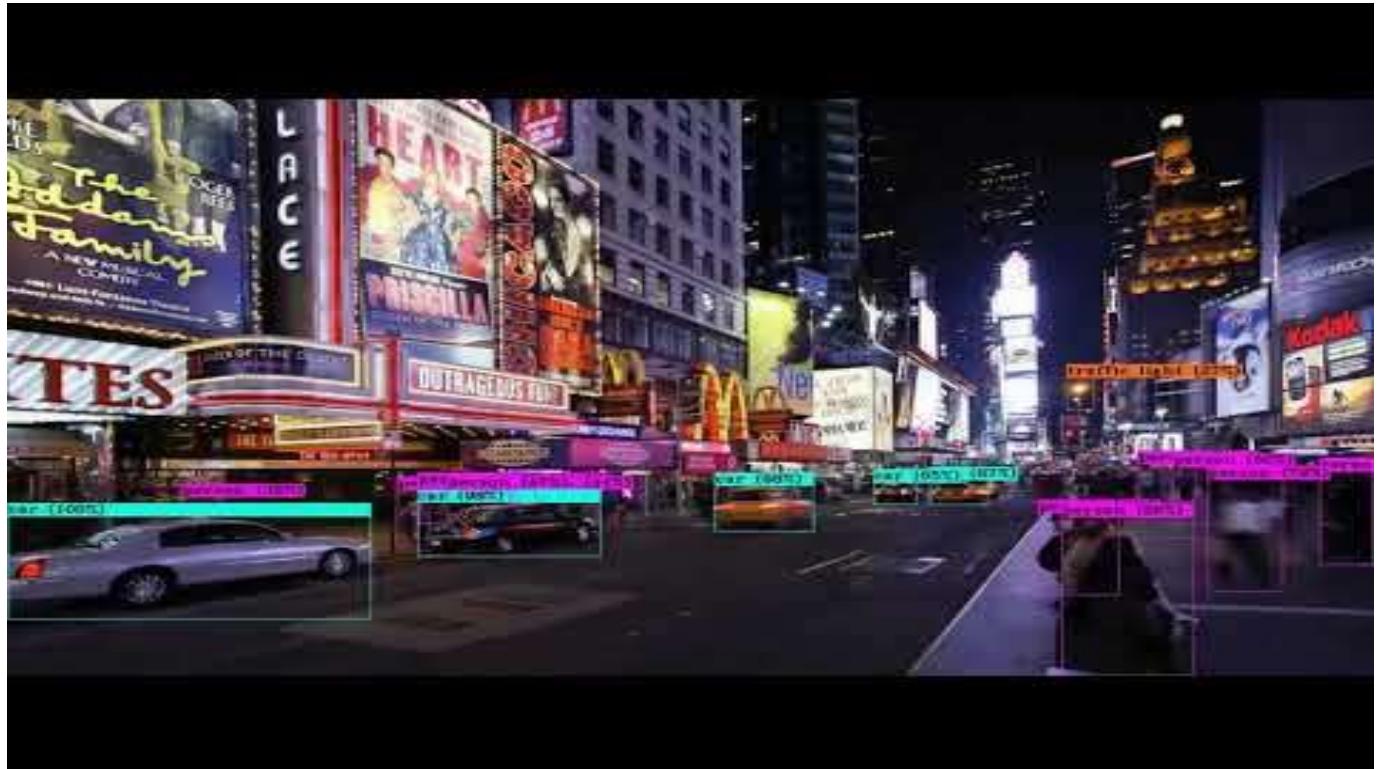


# Intersection Over Union (IoU)

- a single object can have **multiple grid box candidates** for prediction
- **IoU** (between 0 and 1) helps to discard grid boxes to keep only the relevant ones
  - IoU computed for each grid cell = Intersection area divided by the Union Area.
  - IoU selection **threshold** is defined parameter (0.5)
  - ignore the prediction of the grid cells having an  $\text{IOU} \leq \text{threshold}$



Result



[https://www.youtube.com/watch?v=1\\_SiUOYUoOI](https://www.youtube.com/watch?v=1_SiUOYUoOI)

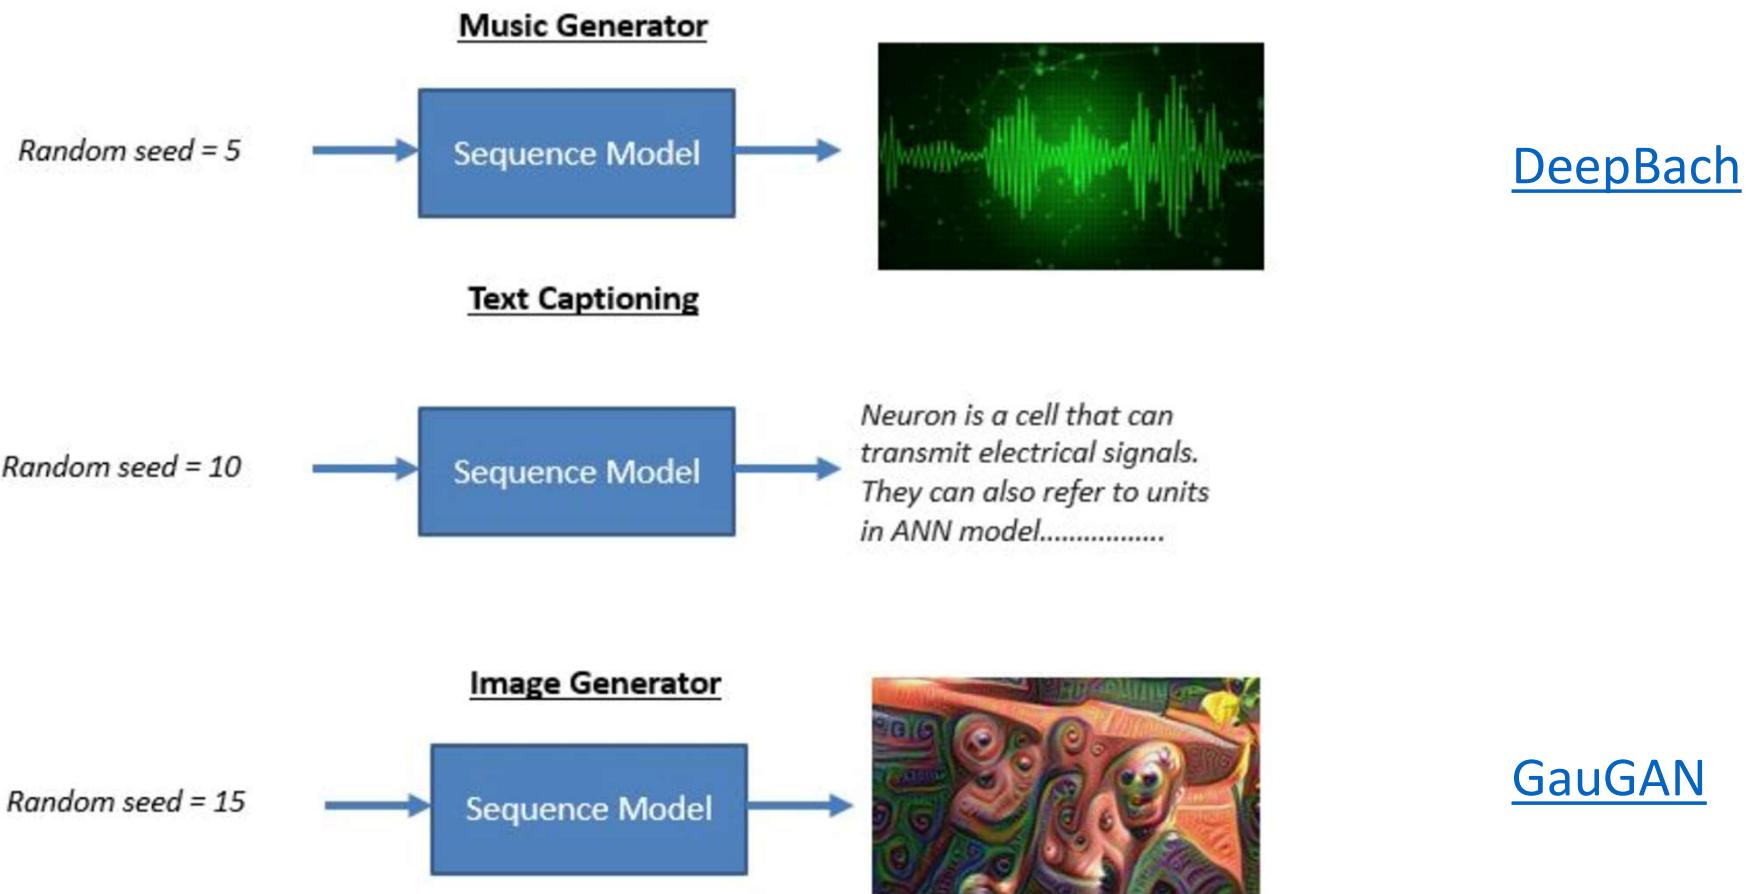
# Generative Models

Holy grail of Deep Learning  
these days

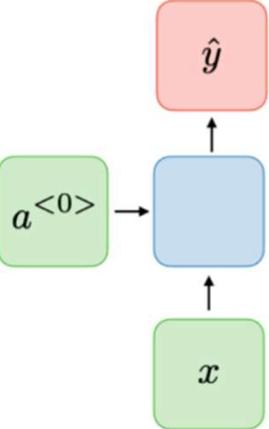


DALL-E : 'A photograph of a cow on the moon.'

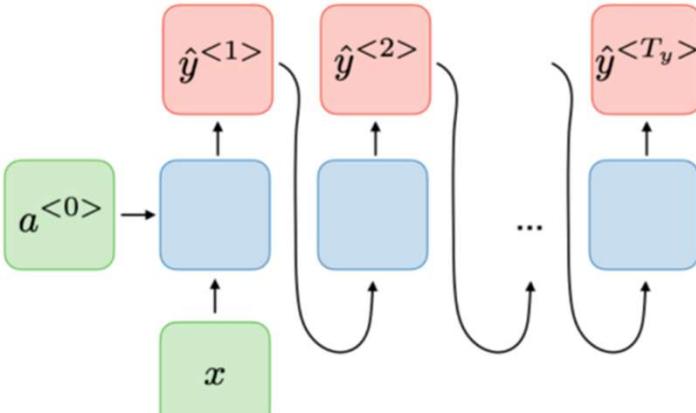
# Generative Models Examples



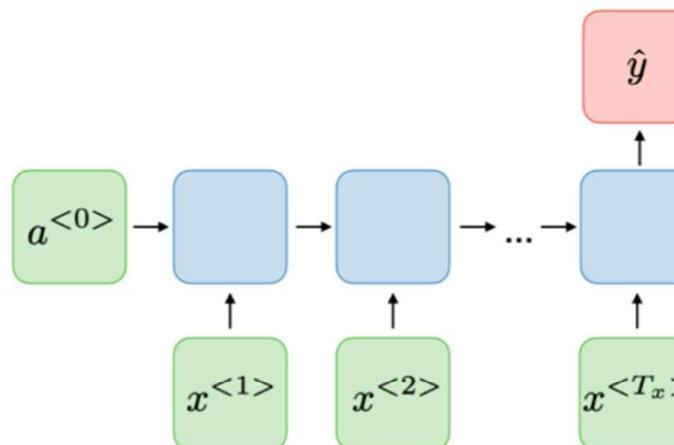
# Applications of RNN (Recurrent Neural Network)

Type of RNN	Illustration	Example
One-to-one $T_x = T_y = 1$	 <pre>graph TD; x[x] --&gt; a0[a&lt;0&gt;]; a0 --&gt; y_hat[y-hat]</pre>	Traditional neural network

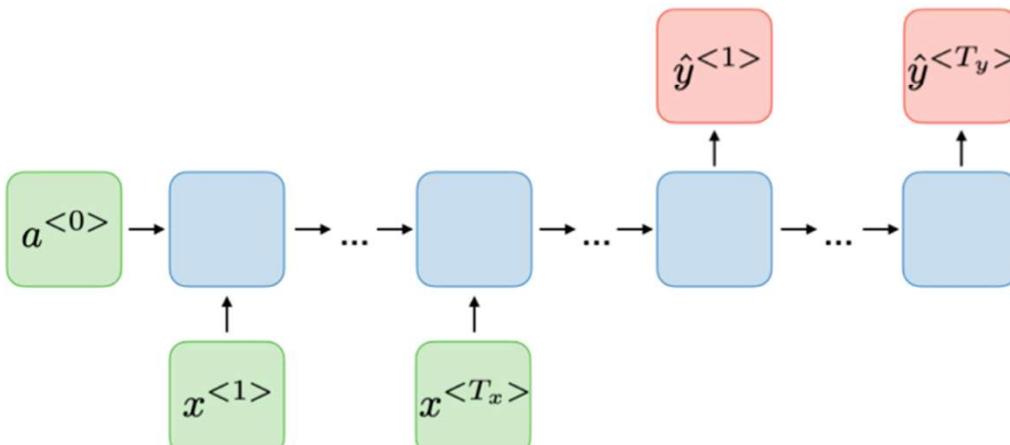
# Applications of RNN

Type of RNN	Illustration	Example
One-to-many $T_x = 1, T_y > 1$		Music generation

# Applications of RNN

Type of RNN	Illustration	Example
Many-to-one $T_x > 1, T_y = 1$	 <p>The diagram illustrates a Many-to-one Recurrent Neural Network (RNN) architecture. It starts with an initial hidden state <math>a^{&lt;0&gt;}</math> (green box). This state is processed by a sequence of recurrent units (blue boxes), each receiving an input <math>x^{&lt;1&gt;} \rightarrow x^{&lt;2&gt;} \rightarrow \dots \rightarrow x^{&lt;T_x&gt;}</math> (green boxes) and producing a hidden state <math>h^{&lt;1&gt;} \rightarrow h^{&lt;2&gt;} \rightarrow \dots \rightarrow h^{&lt;T_x&gt;}</math>. The final hidden state <math>h^{&lt;T_x&gt;}</math> is passed through an output layer (red box) to produce the predicted output <math>\hat{y}</math>.</p>	Sentiment classification

# Applications of RNN

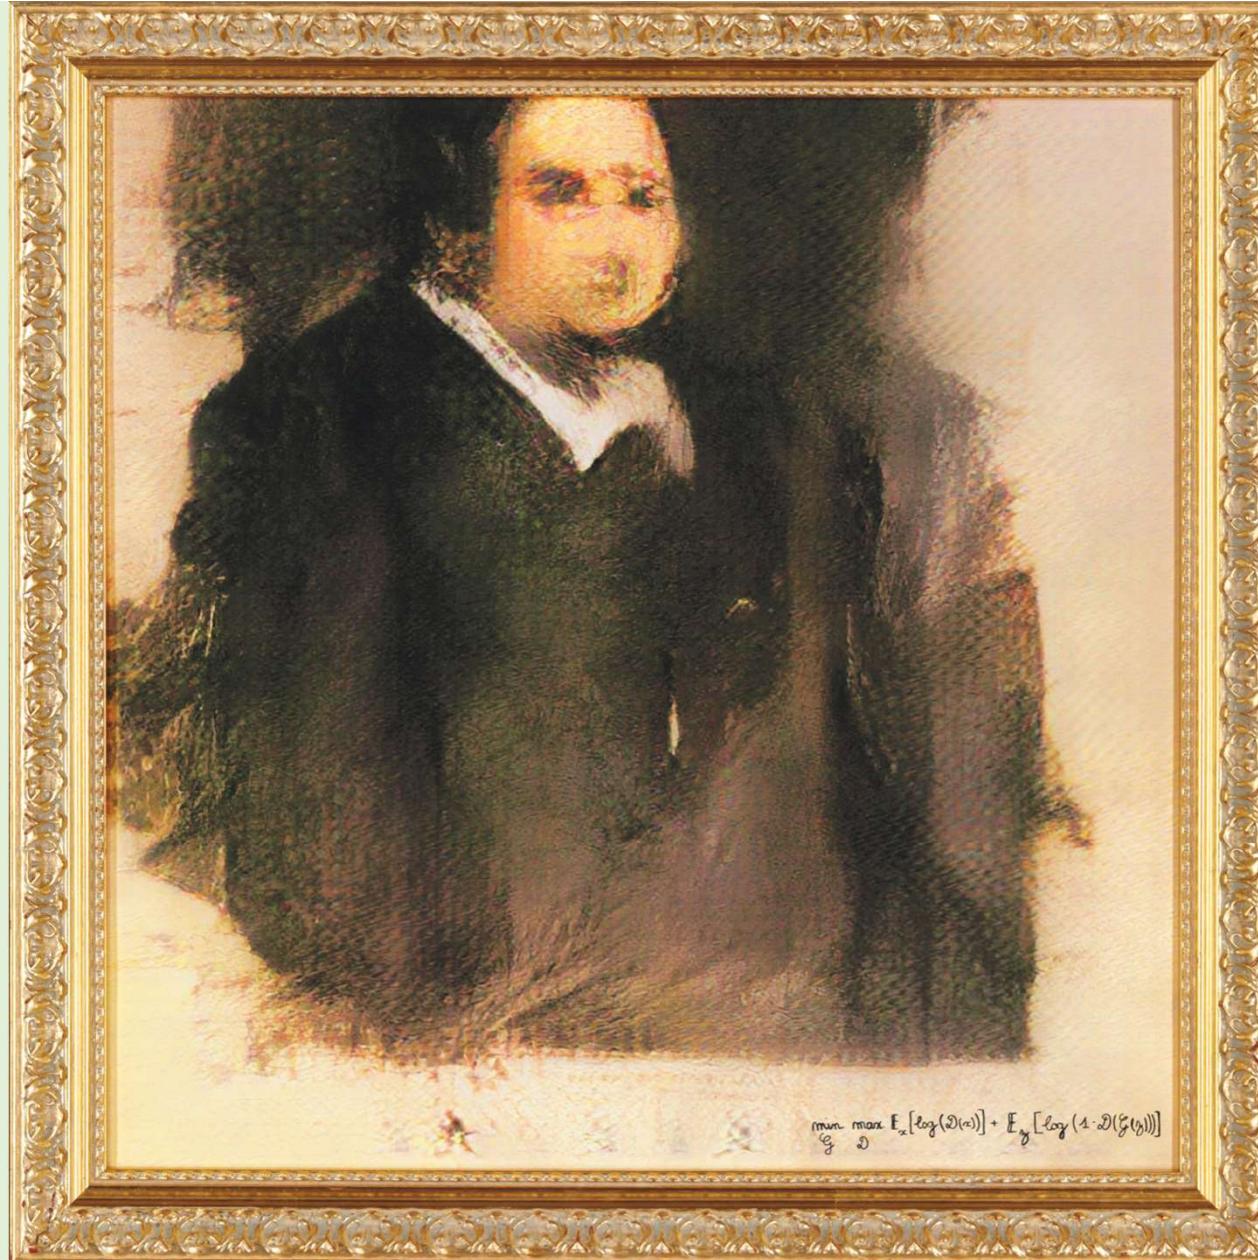
Type of RNN	Illustration	Example
Many-to-many $T_x \neq T_y$		Machine translation

# Generative Adversarial Network

Invented by Ian Goodfellow

*How much are you  
ready to pay for it ?*

Christie's New York 2018



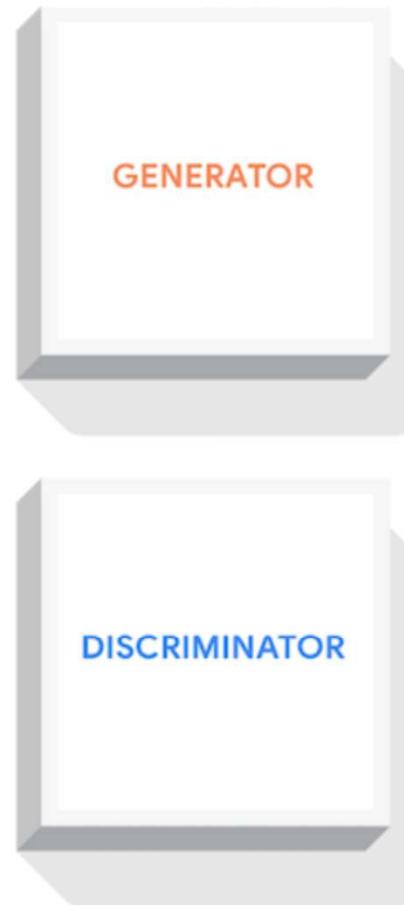
# Principle

**GENERATOR**  
“The Artist”  
A neural network trying to  
create pictures of cats that  
look real.

→

**DISCRIMINATOR**  
“The Art Critic”  
A neural network examining  
cat pictures to determine if  
they’re real or fake.

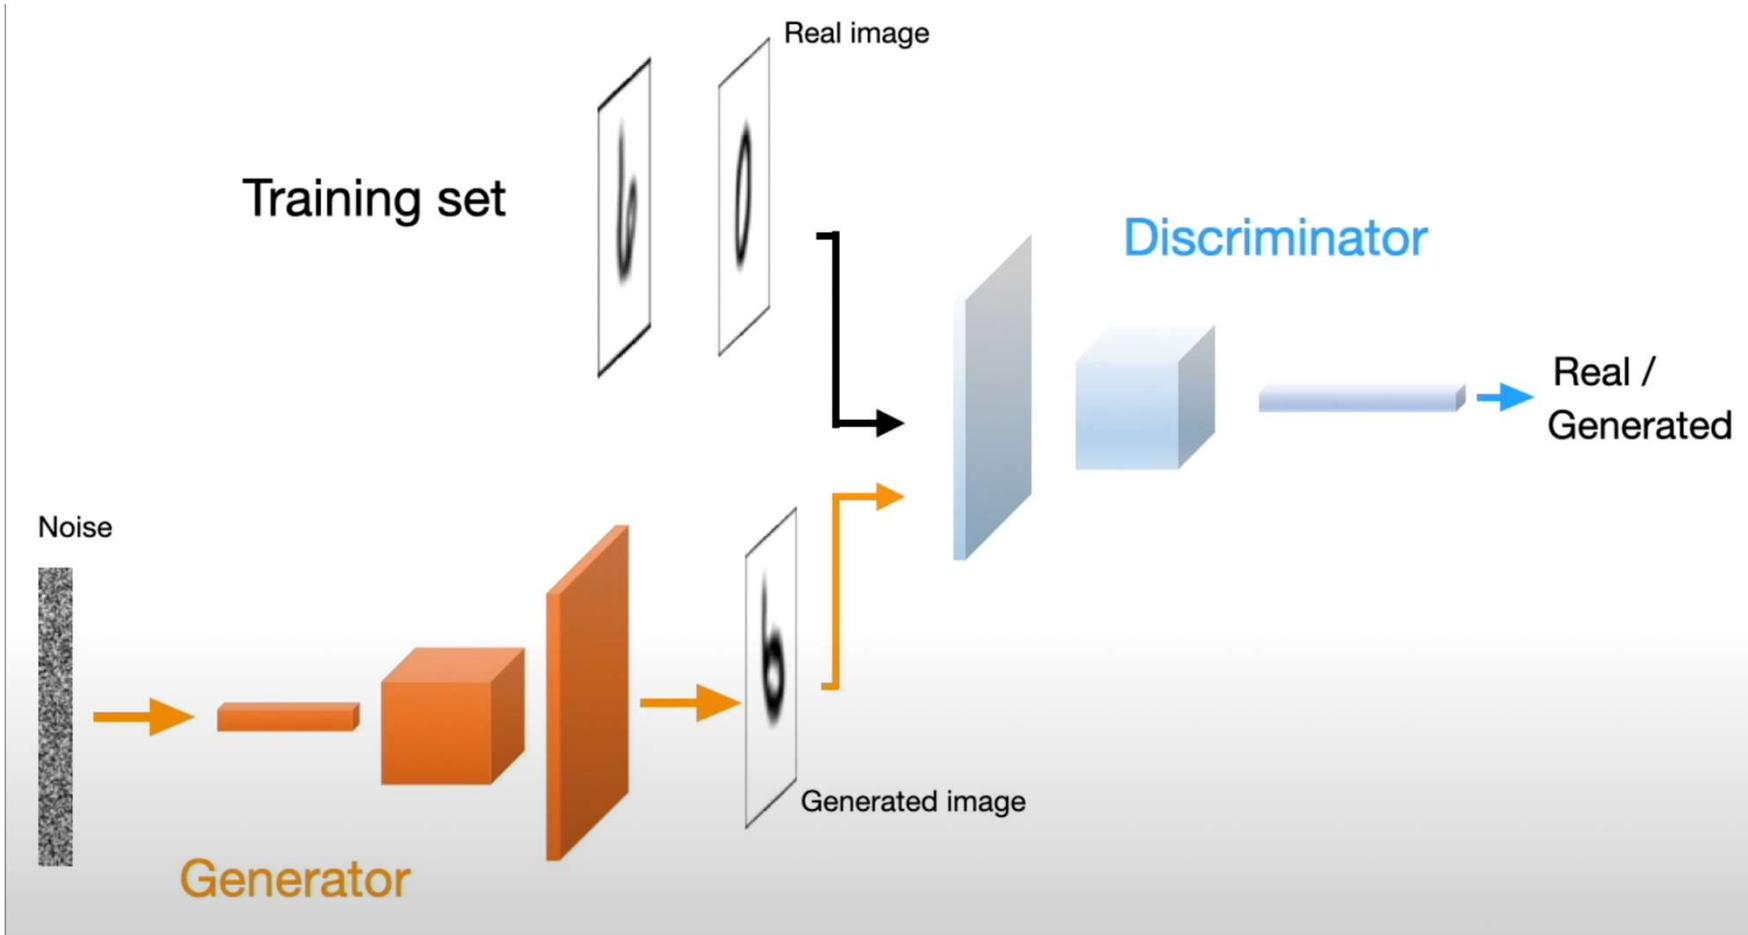
→



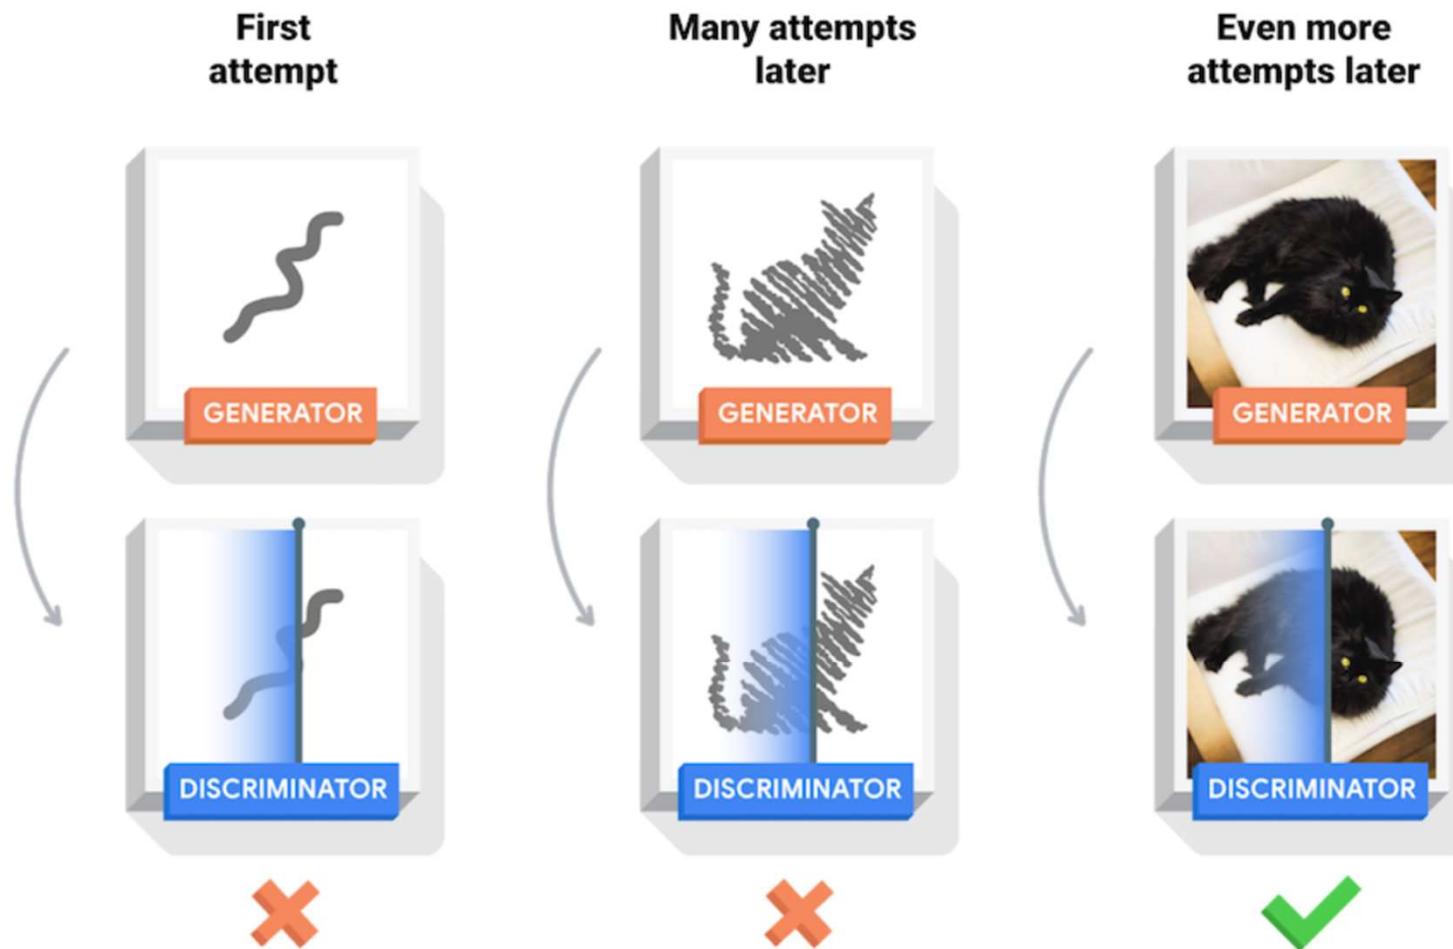
Thousands of real-world  
images labeled “CAT”



# Principle



# Principle



# GAN use case : generate images



Figure 3. Example results by our proposed StackGAN, GAWWN [20], and GAN-INT-CLS [22] conditioned on text descriptions from CUB test set. GAWWN and GAN-INT-CLS generate 16 images for each text description, respectively. We select the best one for each of them to compare with our StackGAN.

# State Of The Art in GANs



(Karras et al, 2018)

(Brock et al, 2018)

# Tutorial / Practical

